

2021 Final - MCQ

Question 1

We want to return, from the two posting lists below, the top-2 documents matching a query using Fagin's algorithm with the aggregation function taken as the sum of the tf-idf weights. How many entries (total of both lists) are accessed in the first phase of the algorithm performing round robin starting at List 1 (i.e., before performing the random access)?

List 1

Document	tf-idf
d3	0.8
d2	0.6
d1	0.5
d4	0.4

List 2

Document	tf-idf
d1	0.8
d3	0.6
d4	0.5
d2	0.4

- a. 4
- b. 8
- c. 6 (CORRECT)
- d. 2

Question 2

Which of the following is **true** regarding inverted files?

- a. Storing differences among word addresses reduces the size of the postings file (CORRECT)
- b. Varying length compression is used to reduce the size of the index file
- c. The space requirement for the postings file is $O(n\beta)$, where β is generally between 0.4 and 0.6

- d. Inverted files prioritize efficiency on insertion over efficiency on search

Question 3

The number of non-zero entries in a column of a term-document matrix indicates:

- a. none of the other responses is correct
- b. how relevant a term is for a document
- c. how often a term of the vocabulary occurs in a document
- d. how many terms of the vocabulary a document contains (CORRECT)

Question 4

Which of the following is **true** for community detection in social graphs?

- a. If n cliques of the same order are connected cyclically with n edges, then the Louvain algorithm will always detect the same communities, independently of the order of processing of the nodes. (CORRECT)
- b. The Louvain algorithm always creates a hierarchy of communities with a common root.
- c. The Louvain algorithm is efficient for small networks, while the Girvan-Newman algorithm is efficient for large networks.
- d. The result of the Girvan-Newman algorithm can depend on the order of processing of nodes whereas for the Louvain algorithm this is not the case.

Question 5

When using matrix factorization for information extraction the entries of the matrix are obtained

- a. from both text and a knowledge base (CORRECT)
- b. from a knowledge base represented as text
- c. from text
- d. from a knowledge base

Question 6

For the number of times the apriori algorithm and the FPgrowth algorithm for association rule mining are scanning the transaction database the following is true

- a. fpgrowth and apriori can have the same number of scans (CORRECT)
- b. apriori cannot have fewer scans than fpgrowth
- c. fpgrowth has always strictly fewer scans than apriori
- d. all three above statements are false

Question 7

Why is non-discounted cumulative gain used as evaluation metrics for recommender systems

- a. because it is more accurate than retrieval metrics, like precision and recall
- b. because often only the top recommendations are considered by the user (CORRECT)
- c. because it allows to consider the financial value of recommended items
- d. because it considers the predicted ratings of all items that have not been rated by the user

Question 8

Which of the following models for generating vector representations for text require to precompute the frequency of co-occurrence of words from the vocabulary in the document collection

- a. CBOW
- b. Fasttext
- c. Glove (CORRECT)
- d. LSI